

KNN Algorithm

- KNN is K-Nearrest Neighbor model which is a supervised classification algorithm used to predict discrete class labels/categories. Here the target variable is a categorical value. It could be a binary class or multi class.
- KNN works based on the majority of the K-nearest neighbors/K-nearest observations.
- KNN Algorithm, 1. Picks a value for K. 2. Calculates the distance from the unknown data point to k neighbors. 3. Select all the k-neighbors or the k-observation in the training data set that are nearest to the unknown data point. 4. Predict the response of the unknown data point using the most popular response value from the k-nearest neighbors.

Note: 1. K value is given by the user but the best k-value can be determined by trial and error method of checking the accuracy score. 2. Distance is calculated using the Euclidean's distance formulae.

```
# Importing necessary libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Onboarding data onto colab
```

```
from google.colab import files
rawdata=files.upload()
```

```
<IPython.core.display.HTML object>
```

```
Saving diabetes (1).csv to diabetes (1).csv
```

```
# DataFrame
```

```
df=pd.read_csv('diabetes (1).csv')
df
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..

763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

Shallow copy

df_copy=df.copy()

Exploratory Data Analysis

df.head()

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

DiabetesPedigreeFunction	Age	Outcome
--------------------------	-----	---------

0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
df.tail()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

```
df.shape
```

```
(768, 9)
```

```
# Technical report
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

```
# Statistical repopr
```

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000
75%	6.000000	140.250000	80.000000	32.000000
max	17.000000	199.000000	122.000000	99.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
df.corr()
```

	Pregnancies	Glucose	BloodPressure	
SkinThickness \				
Pregnancies	1.000000	0.129459	0.141282	-
Glucose	0.129459	1.000000	0.152590	
BloodPressure	0.141282	0.152590	1.000000	
SkinThickness	-0.081672	0.057328	0.207371	
Insulin	-0.073535	0.331357	0.088933	
BMI	0.017683	0.221071	0.281805	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	

```

0.183928
Age          0.544341  0.263514    0.239528    -
0.113970
Outcome      0.221898  0.466581    0.065068
0.074752

```

```

          Insulin      BMI  DiabetesPedigreeFunction
\
Pregnancies      -0.073535  0.017683      -0.033523
Glucose           0.331357  0.221071      0.137337
BloodPressure     0.088933  0.281805      0.041265
SkinThickness     0.436783  0.392573      0.183928
Insulin           1.000000  0.197859      0.185071
BMI               0.197859  1.000000      0.140647
DiabetesPedigreeFunction  0.185071  0.140647      1.000000
Age              -0.042163  0.036242      0.033561
Outcome           0.130548  0.292695      0.173844

```

```

          Age      Outcome
Pregnancies      0.544341  0.221898
Glucose           0.263514  0.466581
BloodPressure     0.239528  0.065068
SkinThickness     -0.113970  0.074752
Insulin           -0.042163  0.130548
BMI               0.036242  0.292695
DiabetesPedigreeFunction  0.033561  0.173844
Age              1.000000  0.238356
Outcome           0.238356  1.000000

```

Check for null value

```
df.isna().sum()/len(df)*100
```

```

Pregnancies      0.0
Glucose           0.0
BloodPressure     0.0
SkinThickness     0.0
Insulin           0.0
BMI               0.0
DiabetesPedigreeFunction  0.0
Age              0.0

```

```

Outcome                                0.0
dtype: float64


# check for duplicate value

df.duplicated().sum()

0

sns.heatmap(df.corr(),annot=True,cmap='BuPu')
plt.show()


```



```


plt.scatter(df['Pregnancies'],df['Outcome'],alpha=0.5,color='green')
plt.show()

```



Multi-variate Analysis

```
sns.regplot(x='BMI',y='Age',data=df)  
plt.show()
```



Data cleaning

Since there is no null values or duplicate values, skipping this data cleaning step.

Data encoding

Since all the features have numerical value in it, skipping this data encoding step.

Standardization

```
X=df.iloc[:, :-1]
Y=df.iloc[:, [-1]]

from sklearn.preprocessing import StandardScaler

sc=StandardScaler()
X_sc=sc.fit_transform(X)
X_sc
```



```
array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,
         0.46849198,  1.4259954 ],
       [-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,
        -0.36506078, -0.19067191],
       [ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,
         0.60439732, -0.10558415],
       ...,
       [ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,
        -0.68519336, -0.27575966],
       [-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,
        -0.37110101,  1.17073215],
       [-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,
        -0.47378505, -0.87137393]])
```

Train Test Split

```
from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(X_sc,Y,test_size=0.2,random_state=7)
```

Model Building

```
from sklearn.neighbors import KNeighborsClassifier

knc=KNeighborsClassifier(n_neighbors=7,p=2)
knc.fit(x_train,y_train)

/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:215: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)

KNeighborsClassifier(n_neighbors=7)

y_predict=knc.predict(x_test)

y_predict
array([0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0,
       0,
       1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0,
       0,
       0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
       1,
       0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,
```

```
0,
    1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
1,
    0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1,
    0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0,
0])
```

y_test

	Outcome
353	0
236	1
323	1
98	0
701	1
..	...
153	0
392	0
308	1
70	1
513	0

[154 rows x 1 columns]

Model Evaluation

```
from sklearn.metrics import
confusion_matrix, accuracy_score, recall_score, f1_score

cm=confusion_matrix(y_test,y_predict)
acs=accuracy_score(y_test,y_predict)
rs=recall_score(y_test,y_predict)
f1=f1_score(y_test,y_predict)
print('Confusion_matrix',cm)
print('Accuracy_score',acs)
print("Recall_score",rs)
print('F1_score',f1)

Confusion_matrix [[82 15]
 [29 28]]
Accuracy_score 0.7142857142857143
Recall_score 0.49122807017543857
F1_score 0.5599999999999999
```

F1_score reveals that the model's accuracy or learning capability is 0.5, (i.e) 50% of the data is being predicted correctly by the model.